

DHRUV VASHIST

Data Scientist | LLM Applications | MLOps Specialist

+919882261260 | dhruvvashist73@gmail.com | linkedin.com/in/dhruv-vasishta/

PROFESSIONAL SUMMARY

Data Scientist with 6+ years designing production LLM applications and scalable ML solutions. Expert in FastAPI, PySpark, Kubernetes orchestration, and LLM monitoring. Reduced operational time by 80% and increased revenue by 20%.

TECHNICAL SKILLS

- LLM & GenAI:** Multi-Agent Systems, RAG, Vector Search, LangChain, LangGraph, LangSmith, GPT-4o, LLaMA, Fine-Tuning (QLoRA, PEFT)
- Development:** Python, SQL, PySpark, FastAPI, Flask, TensorFlow, Pandas, NumPy, Scikit-learn, SpaCy
- MLOps:** Kubernetes, Docker, CI/CD, Git, MLFlow, Prometheus, Model Monitoring
- Cloud:** Azure AI Studio, AWS Bedrock, AKS, AWS EKS, SageMaker, CloudWatch

PROFESSIONAL EXPERIENCE

Senior Data Scientist | Sekhmet Technologies (THB)

Sept 2021 – Oct 2025

LLM-Based Agentic AI Chatbot

- Lead team of data scientists to develop a RAG based chatbot by leveraging GPT-4o as the LLM model and deployed it on client's server using FastAPI
- Developed multi-agent system with Responsible, Rewrite, and SQL agents using LangChain and LangGraph
- Implemented Conversational Agents (Chat Zero Shot REACT Description) to maintain chat history
- Integrated Chroma vector database for efficient vector search and retrieval
- Automated output generation in tables, summaries, and visualizations (charts/graphs) via Output Agent
- Reduced time to access critical business insights by 80% through optimized LLM pipeline

Multi-Agent LLM Chatbot for Business Intelligence

- Designed production 6-agent LLM application with Text-to-SQL.
- Deployed on Azure App Services with CI/CD automation via GitHub Actions
- Monitored LLM performance using LangSmith ensuring continuous optimization and 99.9% uptime

Deep Learning Recommendation System

- Optimized TensorFlow recommendation system with FastAPI endpoints serving 196+ pharmacies
- Orchestrated containerized deployment using Docker and AKS with Prometheus monitoring
- Built PySpark pipelines using databricks achieving 20% cost reduction.

Data Scientist | Wipro Limited

Aug 2019 – Sept 2021

Smart Report & Master Data Management

- Developed production NLP pipelines with FastAPI for EHR processing (1.7M+ customers)
- Built scalable PySpark ETL workflows for data transformation with automated monitoring

EDUCATION

B.Tech in Electrical and Electronics Engineering | NIT Hamirpur

2015 – 2019

CO-CURRICULAR

- Captain of College Tennis Team in several Inter-NIT Tournaments.
- Passionately follows football (Hala Madrid !).